


 BM40A1500 DATA STRUCTURES AND ALGORITHMS

BALANCED TREES AND HEAPS

2024



THE PROBLEM WITH BINARY SEARCH TREES

- » The search operation in BST is efficient only if the tree is balanced.
- » Depending on the order in which the keys are added, the tree might end up unbalanced.



2-3 TREE

- » A node contains one or two keys.
- » Every internal node has either two children (if it contains one key) or three children (if it contains two keys).
 - » Left subtree contains keys that are smaller than both keys in the root node.
 - » Possible middle tree contains keys that are between the the two key values in the root node.
 - » Right subtree contains keys that are larger than both keys in the root node.
- » All leaves are at the same level in the tree
→ the tree is always height balanced.






2-3 TREE

» If after insertion, a node contains three keys, transformations are done to satisfy the criteria.

» Split-and-promote process:


- » The middle key is “promoted” to the parent node and
- » The original node is split into two.
- » Repeated recursively if needed.




AVL TREE

- A self-balancing BST.
- Balance property: for every node, the heights of its left and right subtrees differ by at most 1.
- To maintain the balance property during insertion (and delete), rotations are done when needed:
 - ❖ Single rotation
 - ❖ Double rotation


Balanced




Imbalanced




Case 1: a single rotation



Case 2: a double rotation needed



Still imbalanced



Symmetrical cases similarly

IMPLEMENTING AVL TREE IN PYTHON

```
class AVLNode:
    # Initialize new node
    def __init__(self, key):
        self.key = key
        self.left = self.right = None
        self.balance = 0

class AVL:
    # Initialize new tree
    def __init__(self) -> None:
        self.root = None
        self.is_balanced = True

    # Inserts a new key to the search tree
    def insert(self, key: int):
        self.root = self.insert_help(self.root, key)
```

```
def insert_help(self, root, key):
    if not root:
        root = AVLNode(key)
        self.is_balanced = False
    elif key < root.key:
        root.left = self.insert_help(root.left, key)
        if not self.is_balanced: # Check for possible rotations
            if root.balance >= 0: # No Rotations needed
                self.is_balanced = root.balance == 1
                root.balance -= 1
            else: # Rotation(s) needed
                if root.left.balance == -1:
                    root = self.right_rotation(root) # Single
                else:
                    root = self.left_right_rotation(root) # Double
                self.is_balanced = True
    elif key > root.key:
        root.right = self.insert_help(root.right, key)
    # ...

    return root
```

IMPLEMENTING AVL TREE IN PYTHON

```
# Single rotation: right rotation around root
def right_rotation(self, root):
    child = root.left                # Set variable for child node
    root.left = child.right          # Rotate
    child.right = root
    child.balance = root.balance = 0 # Fix balance variables
    return child

# Double rotation: left rotation around child node followed by right rotation around root
def left_right_rotation(self, root: AVLNode):
    child = root.left
    grandchild = child.right         # Set variables for child node and grandchild node
    child.right = grandchild.left     # Rotate
    grandchild.left = child
    root.left = grandchild.right
    grandchild.right = root
    root.balance = child.balance = 0 # Fix balance variables
    if grandchild.balance == -1:
        root.balance = 1
    elif grandchild.balance == 1:
        child.balance = -1
    grandchild.balance = 0
    return grandchild
```


RED-BLACK TREE


» Each node stores the color: red or black

- used to ensure that the tree stays balanced

» Requirements:

- Each node is either red or black
- Empty sub trees (NIL) are black.
- Red nodes do not have red child nodes.
- Each path from a root to a leaf node contains the same number of black nodes.
- The root node is black.

» To maintain the requirements during the insert and remove operations, color changes and rotations are done if necessary.



COMPARISON TO HASH TABLES

» Hash tables:



- » The basic operations (insert, remove, search) are very fast when the hash table is properly implemented.
 - » Average case: $\Theta(1)$
 - » Worst case: $\Theta(n)$
- » Easier to implement
- » Suitable only for exact-match queries (e.g., searching based on a key value)

» Balanced trees

- » The basic operations are slower than in hash tables
 - » Average case: $\Theta(\log n)$
 - » Worst case: $\Theta(\log n)$
- » Keep data sorted → allow more complex queries
 - » Range queries (finding all records with key value between A and B)
 - » Sequential access
- » Dynamic tree structures are more memory efficient (a hash table needs to be large to be efficient)

HEAPS

- » Simplified version of binary tree.
- » Heap condition:
 - Max heap: every node stores a value that is greater than or equal to the value of either of its children.
 - Min heap: every node stores a value that is less than or equal to that of its children.
- » Efficient to access and remove the smallest (or largest) key and, to add keys.
 - » Commonly used as an implementation for priority queue.
- » Can be implemented as an array:
 - » Tree structure is always complete.



HEAPS

$$\Theta(\log n)$$


» Insert:

- » Insert the new value as a leaf by keeping the tree structure complete.
 - » Put the new value at the end of the array.
- » Move the value upwards on tree until the heap condition is satisfied.

» Remove

- » Swap the root and last leaf.
 - » Swap the first and last position in the array.
- » Remove the last leaf.
- » Push the top value down the tree until the heap condition is satisfied.

$$\Theta(\log n)$$



HEAPSORT

- »» A heap can be used for sorting a list:
 - »» Building a min heap.
 - »» Pop/remove all the elements one-by-one.
- »» Instead of inserting values one-by-one, a heap can be built from an arbitrary array very efficiently:
 - »» Working from bottom to top, from right to left, do a shiftdown on each internal node.
 - »» $\Theta(n)$ (see the proof in OpenDSA material)
- »» Note: when using the array implementation, the max heap is more convenient as no additional array is needed.

